

1 Introduction

The orchestration of containerized applications has become a critical aspect of modern software deployment strategies, particularly as organizations shift towards using microservices across cloud, edge, and fog computing environments. While Kubernetes has emerged as the de facto standard for container orchestration[1], several other solutions, such as Apache Mesos, Docker Swarm, and Cattle, have also gained traction. However, these tools were primarily designed for centralized environments and lack inherent capabilities for managing applications across geo-distributed infrastructures. This review explores an innovative approach, outlined in the research paper on Ge-Kube[2], which extends Kubernetes to address these limitations, specifically focusing on the challenges of deploying containers in a geo-distributed context.

2 Limitations of Existing Solutions

We have listed the capabilities and limitations of Apache Mesos, Docker Swarm, and Cattle in managing containers across multiple geographies, highlighting the challenges they face and the need for additional tools or customizations to address these shortcomings.

Apache Mesos[3], while supporting multi-region deployments with multiple clusters, faces challenges in achieving consistent resource management, scheduling, and data synchronization across these clusters. Each Mesos cluster operates independently, making it difficult to achieve seamless orchestration across regions without additional tools and custom integrations.

Docker Swarm[4, 5], on the other hand, is primarily designed for single-region deployments within a single cluster. While it supports running multiple Swarm clusters, it lacks native support for multi-region deployments with consistent service discovery, networking, and data synchronization across regions, requiring third-party tools or custom scripts to address these limitations.

Cattle[6], a lightweight, event-driven orchestration platform, is not designed for multi-region deployments out of the box. Its metadata-driven service discovery and networking features are tailored for single deployment environments, making it challenging to manage containers consistently across multiple geographies without additional customizations.

3 Challenges in Geo-Distributed Container Orchestration

The advent of cloud computing has transformed the deployment and scalability of applications. However, as the edge computing paradigm gains prominence, it becomes essential to decentralize applications to bring computation closer to data sources, enhancing responsiveness and reducing latency. Traditional orchestration tools like Kubernetes were not originally designed for this shift, leading to inefficiencies in handling dynamically changing workloads over geographically spread out resources.

Kubernetes operates on a cluster-based model where containers are deployed across several nodes managed as a single entity. This model is suboptimal for geo-distributed environments where data must traverse significant network distances, introducing latency that can degrade application performance. Additionally, Kubernetes' scaling policies are predominantly static and threshold-based, relying heavily on cluster-level metrics that do not account for the nuanced demands of latency-sensitive applications distributed across diverse geographic locations.

4 Ge-Kube Framework

The ge-kube framework proposes a novel solution by integrating elasticity management and network-aware container placement into Kubernetes. The framework's architecture is built on decentralized control loops that dynamically adjust container deployment based on real-time network conditions and application demands. This approach not only addresses the scalability and latency issues but also enhances the overall flexibility of deployment strategies.

Ge-kube employs a model-based reinforcement learning approach to manage the scaling of containers. This method dynamically adjusts the number of container replicas in response to application performance metrics, moving beyond the rigid threshold-based policies used in traditional Kubernetes.

The placement manager within ge-kube utilizes a network-aware heuristic that considers inter-node network delays when determining where to deploy containers. This strategy ensures that containers are strategically placed to minimize latency, particularly critical for performance-sensitive applications.

The effectiveness of ge-kube was tested through a series of experiments involving a CPU-intensive application and the Redis database. Results indicated significant performance improvements in terms of reduced latency and increased transaction throughput compared to Kubernetes' default deployment policies.

5 Conclusion

Ge-kube's approach is distinct from other orchestration solutions such as Docker Swarm and Apache Mesos, which do not inherently support geo-distributed deployment or dynamic scaling based on network conditions. The ge-kube framework not only fills this gap but also offers a modular and adaptable solution that can be tailored to specific Quality of Service (QoS) requirements.

The ge-kube framework marks a significant advancement in the field of container orchestration, particularly for managing applications across multiple geographies. Its ability to dynamically adapt to changing workload demands and network conditions presents a compelling alternative to traditional Kubernetes, especially for organizations looking to optimize the deployment of latency-sensitive applications in a geo-distributed environment. The ongoing development of ge-kube could potentially set new standards for deploying and managing containerized applications at scale.

References

- [1] Emiliano Casalicchio. "Container Orchestration: A Survey". In: *Systems Modeling: Methodologies and Tools*. Ed. by Antonio Puliafito and Kishor S. Trivedi. Cham: Springer International Publishing, 2019, pp. 221–235. ISBN: 978-3-319-92378-9. DOI: 10.1007/978-3-319-92378-9_14. URL: https://doi.org/10.1007/978-3-319-92378-9_14.
- [2] Fabiana Rossi et al. "Geo-distributed efficient deployment of containers with Kubernetes". In: *Computer Communications* 159 (2020), pp. 161–174. ISSN: 0140-3664. DOI: <https://doi.org/10.1016/j.comcom.2020.04.061>. URL: <https://www.sciencedirect.com/science/article/pii/S0140366419317931>.
- [3] Apache Mesos. *Resource Management*. URL: <https://mesos.apache.org/documentation/latest/#resource-management>.
- [4] Nikhil Marathe, Ankita Gandhi, and Jaimeel M Shah. "Docker Swarm and Kubernetes in Cloud Computing Environment". In: *2019 3rd International Conference on Trends in Electronics and Informatics (ICOEI)*. 2019, pp. 179–184. DOI: 10.1109/ICOEI.2019.8862654.
- [5] Marek Moravcik and Martin Kontsek. "Overview of Docker container orchestration tools". In: *2020 18th International Conference on Emerging eLearning Technologies and Applications (ICETA)*. 2020, pp. 475–480. DOI: 10.1109/ICETA51985.2020.9379236.
- [6] Rancher. *Cattle*. URL: <https://github.com/rancher/cattle>.